

School of Electrical Engineering and Computer Science Division of Theoretical Computer Science

LAB F Firewall: Packet filtering with iptables

NAME			KTH USERNA	ME
DATI	≣ ::			
TEACHING ASSISTANT'S NAMI	≣ ::			
AB F PASSED (TA'S SIGNATURE) ::			

Compile date: September 2, 2022

Computer Security

DD2395 / HT2022

Contents

1	Intr	oduction	1
	1.1	References	1
	1.2	Downloading and running the Oracle VirtualBox Virtual Machine	1
	1.3	Preparation Questions	2
2	Gett	ting Started	4
	2.1	Connecting to the Virtual Hosts on the VM	4
	2.2	Setting up Interfaces	5
	2.3	Setting up Routing	5
3	iptal	bles: Building a Firewall	7
	3.1	Ping and the Internet Control Message Protocol (ICMP)	7
	3.2	Logging and Limits	9
	3.3	Building a firewall	9
4	nma	p: Detecting Server Capabilities	15
	4.1	Enumeration	15
	4.2	Service Discovery	16
	4.3	OS discovery	17
5	Hist	ory	18
A	Lab	Network Map	19
В	FAQ	2 – Common Problems and Fixes	19
C	Flov	v Diagram of iptables	19
D	Con	nmand Line Basics – in case you are not familiar with the shell	21
	D.1	Manual Pager	21
	D.2	Viewing a File	21
	D.3	Editing a File	21
	D.4	Background Jobs	21
	D.5	Clipboard – Copy+Paste	21

1 Introduction

During this lab, you will learn how to manage a firewall, i.e., packet filtering in GNU/Linux using iptables. Moreover, you will learn how to use the network mapper tool nmap to explore systems on the Internet. The lab is also a refresher of your GNU/Linux and computer networking skills.

iptables (also known as netfilter) is the user-space command-line program used to configure the Linux 2.4.x and later packet filtering ruleset, a successor of ipchains. Both netfilter and ipchains were started in 1999 by Paul Russell, a free software developer.



N Deadlines: Preparation questions before coming to the lab, while tasks by the end of the session

Answer the question in Section 1.3 before coming to the lab session. You should have finished the rest of the tasks by the end of the lab help sessions. Therefore, we encourage you to come prepared.

1.1 References

Before and during the lab, you might want to consult some of the references that you will find below.

Your primary reference for GNU/Linux system utilities should be the reference manuals accessed with the manual pager man. You can open a reference manual for a program by issuing man PROGRAMNAME in the command line. Once a manual is open, you can press h to learn how to search and navigate the manual, and press q to quit. Using the manual is pretty much identical to using less. For this lab, we recommend that your look at:

```
man iptables
man iptables-extensions
man nmap
```

These reference manuals can also be found on the Internet:

```
https://linux.die.net/man/8/iptables (incl. extensions)
https://linux.die.net/man/1/nmap
```

As an **introduction**, the GNU/Linux reference manuals are often too dense, so we suggest you start by reading online tutorials such as

```
https://help.ubuntu.com/community/IptablesHowTo
https://www.frozentux.net/iptables-tutorial/iptables-tutorial.html
https://nmap.org/book/port-scanning-tutorial.html
```

Online tutorials are, however, not always up-to-date, correct, or even existant, so you should learn how to rely on your system's reference manuals.

If you are not used to working with a GNU/Linux terminal, have a look at Appendix D or https://ubuntu.com/ tutorials/command-line-for-beginners.

Downloading and running the Oracle VirtualBox Virtual Machine

You will carry out this lab inside a virtual machine; see the canvas page The Dasak VM for up-to-date information. We recommend that you start setting up the VM now and continue reading until the setup process finishes.

1.3 Preparation Questions

Please answer the following questions before coming to the lab session. You can answer all questions using the references mentioned in Section 1.1.

1.3.1 iptables

chains. In this lab, we will use three standard chains: INPUT, FORWARD, and OUTPUT. Explain which packets will pass through each of these chains:
INPUT:
FORWARD:
OUTPUT:
Operators To change a chain, you need to use an operator. Look up the short form command for the following operators and explain what they do:
Append:

Policy:

Filters Filters are used to choose which packets match a rule. Each filter is used to create matches. Some matches requires modules to be loaded with -m MODULENAME to be available. Explain the following matches and how they

-p:
-s:
-d:
-i:
-o:
--sport:
--dport:

can be used:

Insert:

Delete:

List:

Flush:

Jump Targets We use jump targets to decide what to do with a packet once it has matched a rule. What does the ACCEPT target do?
The DROP target?
The REJECT target?
The NFLOG target?
1.3.2 nmap
The nmap tool provides different scan methods to discover and analyze remote hosts. Each of these scan methods uses a specific network protocol. Each network protocol operates on a particular OSI-model layer (e.g.,ARP or OSI layer 2, IP on OSI layer 3, etc.). So we can group the nmap scan methods according to OSI layers.
Name some protocols for OSI layer mentioned below and describe which protocol nmap uses to scan the corresponding OSI layer (name of the scan method, what command/arguments starts it, what is the purpose of the scan method, what are the limitations of the scan method).
OSI layer 3:
OSI layer 4:
OSI layer 7:
Milestone Milestone 1: Report your progress to a lab assistant.

2 Getting Started



Saving time: Work in parallel and ask questions (using the queue system "Stay A While")

To complete the lab during the allocated time, you can do more than one task at a time as you do not have to do all of them sequentially. If you have to wait for a program to finish or a milestone to be signed, consider continuing with another part or reading the details. If you get stuck somewhere, don't hesitate to ask for help (using the queuing system "Stay A While", choosing the Dasak queue: http://queue.csc.kth.se/Queue/Dasak).

The laboratory system consists of a virtualized system with three virtual hosts running in the The Dasak VM and simulating two interconnected networks (see Appendix A). These hosts are running GNU/Linux (Ubuntu 20.04 LTS). On the course web page, your groups are assigned with numbers. These group numbers determine the IP addresses for your networks according to the network topology map in Appendix A. Please, note that the virtual hosts do not have access to the general Internet.

Group Number, Letter and Login Credentials					
Please determine and fill in the following.					
group number #:		(used in the IP addresses)			
inside subnet:		(simulated private network)			
outside subnet:		(simulated public network)			
ě ,		firewall and outside-host)			
username:	student				
passphrase:	time2work				

2.1 Connecting to the Virtual Hosts on the VM

Downloading, setting up, and starting the The Dasak VM is explained on a corresponding Canvas page. After doing this, login to the VM using the username and passphrase provided in the box above. You will see a Ubuntu desktop, where you will only need three terminal windows – one for each virtual host. Open them and type each of the following commands in one of the windows:

```
sudo lxc-start inside-host && sudo lxc-attach inside-host
sudo lxc-start firewall && sudo lxc-attach firewall
sudo lxc-start outside-host && sudo lxc-attach outside-host
```

This boots the containers and gives you a root shell for them.

Now, open an additional terminal and type the following command in it: sudo lxc-start server && exit

This command starts another container needed later in the lab and then exits.

Tips for working with the virtual host terminals

- Use three terminals for the three different virtual hosts.
- After starting a host using the commands above, it is possible to open additional shells using sudo lxc-attach {hostname}. This can be useful for just displaying the log, for example. But don't open too many because this can confuse more than help.
- You may maximize the VM window and resize the VM's terminal windows to get a good overview.
- For shutting down a virtual host at the end of the lab, you can issue shutdown -h now in one of the consoles. But look at the prompt and make sure that you don't shut down the VM and possibly lose progress. Either way, remember that the network interface and iptables configuration we discuss in the following is not persistent and therefore lost after shutting down.

2.2 Setting up Interfaces

Next, you need to configure your network interfaces. You do this by manually assigning them an IP address and a netmask. The network map in the appendix describes which addresses to use on which of the virtual hosts in your VM. Remember to replace # with your group number in all IP addresses (for example, if your # is 4, then the host inside-host has the IP address 192.168.4.2). Use the ip tool ad make sure to apply the configuration on the correct interfaces (eth0, eth1) on each machine.

For example, to assign the IP address 10.18.0.33 with a 20-bit netmask to the interface eth0, issue the following command:

```
ip addr add 10.18.0.33/20 dev eth0
```

Check the configuration by using ip addr without any additional arguments.

2.3 Setting up Routing

Normal hosts do not route packets, and so packet forwarding is disabled by default on a newly installed machine. To ensure that data is routed through the firewall (firewall), one would have to enable traffic forwarding. Otherwise, it is only possible to communicate between directly connected hosts (e. g.,inside-host \leftrightarrow firewall and firewall \leftrightarrow outside-host) but *not* between the two end hosts (e. g.,inside-host \leftrightarrow outside-host).

First, verify that packet forwarding is enabled on the firewall by checking that net.ipv4.ip_forward is "1": sysctl net.ipv4.ip_forward

If it is not, set it to 1 to enable forwarding (only on the firewall host firewall):

```
sysctl -w net.ipv4.ip_forward=1
```

For completeness, you could disable packet forwarding on inside-host and outside-host using: sysctl -w net.ipv4.ip_forward=0

Either way, the end hosts inside-host (internal) and outside-host (external) have to learn how to reach each other. To configure them, use the following command only on the end hosts:

```
ip route add default via ADDRESS
```

to set their default route to the central firewall host firewall. ADDRESS has to be the corresponding address of the directly connected interface on the firewall host.

Do not forget to check the routing table with the ip route command without any additional arguments (be patient, the output may be delayed a bit).

6



Milestone

Verify that you can ping between the two end hosts of your network (inside-host ↔ outside-host).

iptables: Building a Firewall

Note that all the iptables commands have to be run only on the firewall (firewall)!

To view the iptables manual, run the command:

```
man iptables
and to view the manual for iptables extensions
man iptables-extensions
```

More references are listed in Section 1.1.

To list the current state (list of rules and policies for all chains), run the command: iptables -vL

🔗 Saving & presenting your work

In case a virtual host fails, you might lose the rules you have created. Therefore keep a record of what you are doing (on paper or in a file on your local machine). At any time you can use the iptables-save command, to print out the rules you created so far. Note that saving them to a file on the virtual host might not be a good idea, as the complete host can fail. Please see the following link for examples on how to save and restore iptable rules using files.

```
https://www.cyberciti.biz/faq/how-to-save-restore-iptables-firewall-config-ubuntu/
```

You can also use the snapshot feature of the VM to save the state of the virtual machine at different milestones of the assignment. Please see the following link for more information on how to take snapshots using virtualbox.

```
https://www.techrepublic.com/article/how-to-use-snapshots-in-virtualbox/
```

Important points to note for presentation: Please note that if you are working on your lab outside lab hours and want to present the lab at one go during lab sessions, you have to save your work using at least one of these mechanisms stated above for reproducibility during presentation.

Ping and the Internet Control Message Protocol (ICMP)

The ping tool is usually used to test the reachability of a host that implements the Internet protocol (IP). It operates by sending packets using the Internet Control Message Protocol (ICMP) of the type 'echo request' to the host whose reachability is to be tested, and processing the response from that host (if any).

In Listing 1, you can see a sample output of this utility testing the web server of KTH (ping www.kth.se). The results show three ICMP requests (which were successful as the time implies that there was a response) and a statistical summary of the response packets received.

```
PING www.kth.se (130.237.32.143) 56(84) bytes of data.
64 bytes from 130.237.32.143: icmp_req=1 ttl=254 time=10.4 ms
64 bytes from 130.237.32.143: icmp_req=2 ttl=254 time=10.4 ms
64 bytes from 130.237.32.143: icmp_req=3 ttl=254 time=9.31 ms
--- www.kth.se ping statistics ---
3 packets transmitted, 3 received, 0% packet loss, time 2002ms
rtt min/avg/max/mdev = 9.315/10.064/10.478/0.543 ms
```

Listing 1: Sample output in the command line of the ping utility

3.1.1 Blocking ICMP requests

ping the internal host inside-host from the external host outside-host. What happens? Now execute the following command on the firewall host:

```
iptables -A FORWARD -p icmp --icmp-type echo-request -j DROP
```

Now ping the host again. What happens and why?

Now check the output from iptables -vL, what has changed?

To remove the rule from the chain, first list the rules again:

iptables --line-numbers -L FORWARD

There you can find the line number of the rule you just added, and then remove it with:

iptables -D FORWARD line_number

3.1.2 Rejecting ICMP Requests

Now, create a new rule which REJECTS all ICMP echo-requests from the external network to your internal network. Note that iptables rules only match if all the conditions in the rule are true.

n Specifying Source and Destination

For simplicity, we have only one internal and one external host in the lab network setup. However, when specifying the source and destination for iptables rules, try to be more general than specifying only single IP addresses as source or destination. You can accomplish this either by using an IP address range (e. g., 10. #.0.0/20) for the -s and -d option, or by specifying the involved interface with -i and -o.

Now, verify that

- you can ping from the internal host to the external,
- the external host cannot ping the internal, and
- the firewall can ping both hosts.

Can you ping from the external host to the **internal interface** eth0 on the **firewall host** (note: this question is often misunderstood, please read it carefully)? Why/why not?

Can this have any security implications?

What is the difference between rejecting and dropping traffic?

What are the benefits of rejecting, and what are the benefits of dropping?

3.2 **Logging and Limits**

One important task for a firewall is to log rejected or dropped packets, making it easier to trace attacks. We can create a rule with the jump target NFLOG to save information to the system's ulog. (We use NFLOG instead of LOG, since our virtualized system has some constraints and so iptables cannot log directly to syslog.)

Create a rule that logs all rejected ping messages with the string Ping rejected by <your name>: in the log message. How do you create this rule/s?

Make sure the string "Ping rejected by <your name>:" is written in the log message. Read the ulog to check that we save the log messages:

less +F /var/log/ulog/syslogemu.log

Notice that you cannot use tail -f for displaying the ulog. Furthermore, less +F can be quit by first pressing Ctrl+C to interrupt and then pressing q to quit. We can use the module limit to limit how often a rule can be triggered. Use the limit module to ensure that we save no more than five pings each minute.



Milestone

Milestone 2: Report your progress to a lab assistant.

3.3 **Building a firewall**

You will now build a simple firewall for your network. Before you start, make sure that there are no rules left from the previous assignments: flush all rules by running iptables -F (not specifying a chain removes all rules from all chains).



Order of rules

For each step, think about the order of the rules: Is it enough to append a new rule at the end of a chain, or should it go before a certain another rule to work as intended? If you use the -I argument without specifying a number, you place the rule at the top of the chain.

3.3.1 Default Policy

Each chain has a default policy target that details what to do with any packet that does not match any rules present in the chain. We can set the default policy target to ACCEPT or DROP but not to REJECT as it is an extension target. Set all chains to the policy target DROP, and then verify that you cannot send any data through or to the firewall (e.g.,using ping). How do you set the default policies?

If you instead want to REJECT all packages that do not match a rule in the INPUT chain, how would you do it?

We mentioned that REJECT is an extension target. What is an extension target?

3.3.2 Network Permissions

Now it is time to start allowing some carefully chosen traffic through the firewall. Create a rule that allows all traffic originating from the internal network to arrive at the internal interface (eth0) of the firewall host, and a rule that allows all traffic originating from the firewall to reach the internal network. What commands did you use to create these rules?

Make sure you now can reach the internal network from the firewall and the firewall from the internal network.

3.3.3 Permitting a Service

SSH (Secure Shell) is a standard service for remote management of firewalls. Create rules which allow a host from the external network to connect to the firewall with SSH. SSH runs on port 22 and uses only TCP.

Make sure that the ssh daemon *sshd* is running on all hosts by executing the following command: systemctl restart ssh.service

And check the status of the ssh daemon using the following command:

systemctl status ssh.service

Source and destination ports

Note that the well-known port number for service only applies to the machine where the service is running (server), not for devices connected to the service (clients). So if a client machine sends a packet to a specific service on a server machine, the destination port will be the well-known port number for that service, but the source port can be any number. Accordingly, when the server returns an answer packet to the client, the packet will have the well-known port number as a source port, but the client machine will choose the destination port.

Verify that,

- you can connect to the firewall with SSH from the other hosts,
- you cannot connect directly from your external host to the internal host with ssh, and
- you can connect from your external host to the internal host if you first connect to the firewall with ssh and then connect from there with ssh to the internal host (note: think about which username you should use when logging in to the hosts with ssh).

What kind of security advantage does a setup with an SSH terminal server offer?

What kind of security disadvantage does a setup with an SSH terminal server introduce?

See http://www.fail2ban.org, how does the framework fail2ban mitigate this disadvantage?

3.3.4 Stateful Filtering

In most cases, we want to allow hosts on the internal network to connect to the external network, e. g.,the Internet. However, we do not wish for hosts on the Internet to connect to hosts on the internal network. For some protocols, such as TCP this can be done *stateless*, due to the three-way handshake needed to create a connection. By blocking the initial SYN packet in one direction, we can prevent the establishment of connections in that direction while still allowing connections established in the other direction to send packets both ways (as they will never send an initial SYN packet in the direction that we blocked).

However, stateless filtering breaks many protocols; for example, UDP-based protocols cannot easily be allowed in only one direction. Therefore, we need a *stateful* firewall to handle it properly. Some protocols, such as FTP, also break if connections from the Internet are completely denied.

Examine the module state. Use this module to create a rule that allows the hosts on the inside of the firewall to establish connections to the outside. Allow all data that belongs to these connections through the firewall. Keep blocking all other connection attempts from the outside. What command did you use to create this rule?

3.3.5 Opening Ports

Sometimes, you want computers on the outside of the network to access a specific service on the inside. Therefore, your task is to add rules to your firewall so that external computers can reach the *echo* service (port 7) on hosts on the inside both on UDP and TCP. What commands did you use to create these rules?

The echo service (provided by the extended internet service deamon xinetd) is a legacy testing protocol that replies with the same data that is sent to it. This service is disabled by default, but we have activated this on your hosts (you can check this by running lsof -i and look for lines that end in "echo (LISTEN)" or similar).

To test your firewall rules, connect to the echo service that is running on your internal host, from the external host by running telnet IPADDRESS 7 on the external host and observe that all data you send is echoed back to you. Notice that you can quit telnet by pressing Ctrl+] (or Ctrl+5) and then typing quit in the telnet prompt.

3.3.6 Blocking Ports

Sometimes you do not want your *internal* users to connect to the Internet on a specific port at all. One commonly blocked port is 135 used for Windows file sharing¹. Ensure your firewall blocks all traffic on port 135 (both TCP and UDP) from the computers on the internal network. What commands did you use to create these rules?

3.3.7 Verifying Your Setup

Systematically verify your setup, making sure all rules work correctly. To test a specific rule, you can listen to a port with netcat and try to connect to it using telnet.

On the receiving host run:

nc -l PORT

On the sending host run:

telnet IPADDRESS PORT

Furthermore, use the counters in the output of iptables -vL to make sure the packets match the rules you expect. At this point, you should have the following rules active, in this order of priority:

- Connections on **port 135** from the inside hosts are blocked.
- Connections on the **SSH port** are allowed to the firewall host from the outside.
- Connections on the **echo port** are allowed to the internal hosts from the outside.
- Connections **directly to and from the firewall** are allowed from the internal networks.
- Connections from the inside are allowed out.
- Connections from the outside are blocked.



Milestone 3: Report your progress to a lab assistant.

¹Actually there are more ports involved (135-139 on older Windows machines and 445 on more recent ones) but in this lab it is enough to block 135 as one example.

3.3.8 Defending Against SSH Brute-force Attacks

You have allowed connections to the SSH port (hence, its service) of the firewall from any host on the outside network. There is nothing that would avoid a brute-force attack or a Distributed Denial of Service (DDoS) attack because there are no limitations on the amount of times that one particular host can try a username and a password (besides the ones that the SSH server implementation might have internally).

You will practice the knowledge you have acquired about stateful filtering and logging to prevent a host from the outside network to execute a brute force attack.

Let's start by creating a new chain called SSH to handle all SSH traffic. After making the chain, you need to modify the INPUT chain rule that allowed connections to the SSH port on the firewall host from the outside network: send all SSH packages to the 'SSH' chain instead of directly accepting them.

In the 'SSH' chain, first add a rule that accepts all packages from established connections, so the following rules only apply for new connection attempts (and don't slow down legitimate, established connections).

Now, have a look at the documentation of the module recent. It allows you to create a dynamic list of IP addresses for later use, such as counting the number of times a particular IP address tries to start a new connection on a port.

Your next step, with the knowledge of the module recent, is to create the necessary rules in the 'SSH' chain to catch those hosts that are potentially dangerous. We will assume that anyone trying to log in **more than three times in a row within 30 seconds** from the same host is potentially malicious and should be slowed down.

We want not only to block but also to log suspicious packages. But in case of a brute-force attack, we don't wish to flood the log file. So from each suspicious IP address, we will block the fourth, fifth, and all further connection attempts (in 30 seconds), but only log the fourth connection attempt.

Create the first rule so that the source address of the packet is added to a custom dynamic list named SSH_COUNTER Then, add a second rule so that the fifth (and any further) packet attempting to establish an SSH connection and coming from the same source in a time period of 30 seconds gets dropped.

To log suspicious connection attempts, create two rules so that one of the rules logs the fourth SSH connection from the same host address (use the prefix "SSH brute force attacker: "), and the other one blocks the package.

Finally, you have to accept all packets that have not been stopped so far. Why do you need to do this?

At this point, the newly created 'SSH' chain should contain the following rules, in this order of priority:

- A rule accepting all packets from already established connections.
- A rule adding IP addresses to a dynamic list.
- Every fifth (and higher) packet from the same source host address within 30 seconds is dropped.
- Every fourth packet from the same source host address within 30 seconds is logged.
- Every fourth packet from the same source host address within 30 seconds is dropped.
- Anything that has not been stopped so far is accepted.

Test your rules, paying attention to the counters with iptables -vL before and after connecting via ssh. If you implemented the rules correctly, you should notice that you have to wait a few seconds from the moment you start trying to connect for the fourth time in 30 seconds from the external host.



? Visualizing the current list of blocked addresses

The directory /proc/self/net/xt_recent/ stores each dynamic list you define in a file of the same name as the dynamic list's identifier. For testing your rules, you can run cat SSH_COUNTER to print the contents of the file, e. g., before and after each test to see the changes.

You now have a set of rules that would minimize the impact of an attack on your system, but there is still a way for a powerful attacker to carry out a successful attack. Can you speculate how and why?



Milestone

Milestone 4: Report your progress to a lab assistant.

3.3.9 **Building Your Own Firewall**

Consider your home network, the services you offer, and the service you require from the Internet. Which protocols do you use? Which can you block? When you compare these requirements to the rules you just created in this lab, which ones would you keep, which ones not, and which additional rules would you add? You may note them down in abbreviated form (which port/protocol, what to do with it, why).

Rules to keep:

Rules not to keep:

Rules to add



Milestone 5: Report your progress to a lab assistant.

nmap: Detecting Server Capabilities

nmap is a free open source utility commonly used for network exploration and security auditing. nmap can determine what hosts are available on the network, what services (application name and version) those hosts are offering, what operating systems (and OS versions) they are running, what type of packet filters/firewalls are in use, and dozens of other characteristics. To view the nmap help, run: man nmap



No not use nmap outside the lab

In this second part of the lab, you will switch from the perspective of protecting a network with a firewall to the other side and explore ways to attack a system using the network mapper nmap. You can safely explore this tool inside this lab environment. Outside the lab, however, remember never to use nmap on a system if you do not have the explicit consent of the system's owner. Even basic scans on one IP address might be considered as attack and have potentially legal implications.

Enumeration 4.1

You will now use nmap to enumerate your outside network from the external host (outside-host). Try to locate the IP address of the unknown server that is connected to the network (see Appendix A). Note that the scan can take up to 5 minutes, so you might want to test your command on a smaller portion of the network before scanning the whole subnet. Also make sure to use options that speed up the scanning, otherwise the scan will take a very long time (up to one hour with some standard settings). Have a look at the --min-rate option for that purpose.

Saving time

Since the scan will take some time, pipe the output from nmap to a file and put the process in the background so you can continue your work. Use the > operator to pipe the output to a file and the & operator to send the process to background:

```
nmap [arguments for nmap] > [output_file] &
```

In case you have trouble to speed up the search and do not find the server, you may ask an instructor in the lab room for the server IP corresponding to your group number. In this way you can continue with the fingerprinting and service search.

What command did you use to locate the server?

What is the server's address, how long time did the scan take and how many addresses did you scan?

If you get an error message while scanning: Make sure you are root. If you still can't find the server, verify your routing and netmask settings.

4.2 Service Discovery

Now that you have found the server's IP address, try to gather more information, such as the services running on
the server. Scan for both TCP and UDP services. Futhermore, try to find out the version number of the discovered
services. For the UDP scan you might want to look into the -F option.

What command di	d you use for	TCP discovery?
-----------------	---------------	----------------

What command did you use for UDP discovery?

UDP discovery is much slower than TCP discovery, why?

What is the difference between open, closed, filtered and unfiltered ports?

Which services did you find? (If you found about 7 open TCP ports and about 3 open UDP ports you are fine. You might not be able to find version information for all services.)

TCP/UDP	Port	Service	Version

TCP/UDP	Port	Service	Version

4.3 OS discovery

Finally, try to guess the operation system of the server by using the appropriate nmap operations.

What command did you use:

What operation system did nmap consider to be most likely? (If you do not get any text that makes sense to you, try again using the extra option --fuzzy)

What information is used by nmap, to guess the operation system?

Milestone
Milestone 6: Report your progress to a lab assistant.

5 History

Version	Contribution	Author (Affiliation)	Contact
1.0	First development	Pehr Söderman (ICT/KTH)	pehrs@kth.se
2.0	Adaptation for HT2012	Benjamin Greschbach (CSC/KTH)	bgre@kth.se
3.0	Adaptation for HT2013	Guillermo Rodríguez Cano (CSC/KTH)	gurc@csc.kth.se
3.1	New exercise for HT2013	Guillermo Rodríguez Cano (CSC/KTH)	gurc@csc.kth.se
3.2	Updating for HT2015	Benjamin Greschbach (CSC/KTH)	bgre@csc.kth.se
4.0	VirtualBox for HT2017	Andreas Lindner (CSC/KTH)	andili@kth.se
4.1	Adaption for HT2018	Andreas Lindner (EECS/KTH)	andili@kth.se
4.2	VM info pointer for HT2018	Sonja Buchegger (EECS/KTH)	buc@kth.se
4.3	Adaption for HT2019	Md Sakib Nizam Khan (EECS/KTH)	msnkhan@kth.se
4.4	Saving work pointer HT2020	Md Sakib Nizam Khan (EECS/KTH)	msnkhan@kth.se
4.5	Adaption for HT2021	Md Sakib Nizam Khan (EECS/KTH)	msnkhan@kth.se
5.0	Improvements (e.g., text fields)	Henrik Karlsson (EECS/KTH)	henrik10@kth.se
5.1	Adaption for HT2022	Anoud Alshnakat (EECS/KTH)	anoud@kth.se

A Lab Network Map

The network map below shows the connections and information about the addressing of the hosts in the lab network. Remember to replace the placeholder # with your according group value. Furthermore, you have to concatenate the network prefix with the host part to obtain the IP address for a certain interface. The length of the network prefix is given by the network mask². For example, if your group number # is 42, then eth0 of the external host should get the address 10.42.0.2 (10.42.0.0/20 being the network prefix and .0.2 the host part), eth0 of the firewall host should get 192.168.42.1, and so on.

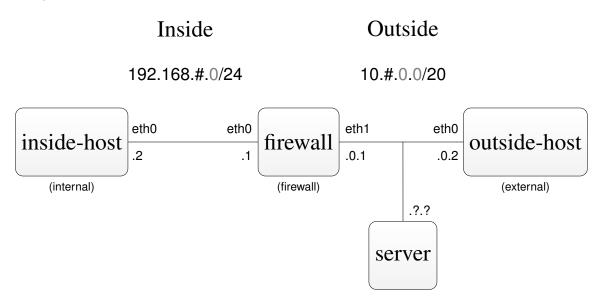


Figure 1: Lab network map.

B FAQ – Common Problems and Fixes

I can't ping between inside-host and outside-host (but between direct neighbors). Check that route add default gw ADDRESS was issued WITHOUT specifying a netmaks after the neighbor's address. You have to take the interface down (ifconfig eth $\frac{0}{1}$ down) and up again (ifconfig eth $\frac{0}{1}$ up), after issuing the correct command.

Telnet got stuck, how can I exit from it? Issue Ctrl+C on the listening netcat no instance. Ctrl+] can be used to exit telnet as usual.

I can't login with ssh when connecting from one virtual host to another. Make sure to use the username student (the only user account on the virtual hosts) when connecting with ssh: ssh student@IPADDRESS. Then the passphrase time2work will work.

I've set up the correct rules, but ping does not work (no rules match). Note that ping uses the ICMP protocol, so it's neither using TCP nor UDP. If your rules match only for TCP and UDP, ping traffic is not affacted by them.

C Flow Diagram of iptables

The following graph shows a simplified and brief scheme of the iptables flow graph. Each chain contains the corresponding tables for the processing of the packets. There are five predefined chains, namely, PREROUTING,

²The netmask is written right behind the network address in standard slash-notation: /24 stands for a netmask of 24 bits (which is equal to 255.255.255.0 in legacy notation). Remember that an IPv4 address has 32 bits, so each number in the dot-notation represents 8 bits.

FORWARD, INPUT, OUTPUT and POSTROUTING, but not all chains have all the tables (e.g., the INPUT chain only contains the 'mangle' and the 'filter' tables). Note that predefined chains have a default policy (e.g., ACCEPT) which is always applied to the packet traversing the chain when it does not match any of the rules defined for that chain.

Packets start at a given chain, but typically they come either as 'inbound traffic' (that is, from the network card) or from a 'local process', and after traversing the tables of each chain, they will get out of iptables as 'outbound traffic'. Note that depending on the rules the processing of the packets might not be sequential, as the graph shows, and the packets may jump to another chain or even discarded/dropped, but in any case, every packet coming or leaving iptables will traverse one chain at least.

For this lab the only table that you will be using is the one focused on packet filtering, 'filter'.

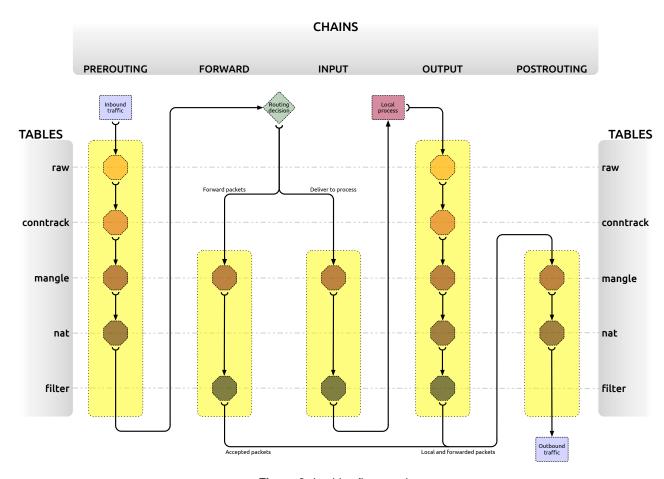


Figure 2: iptables flow graph

D Command Line Basics – in case you are not familiar with the shell

D.1 Manual Pager

You get a reference manual for almost all command line programs by issuing man PROGRAMNAME. It uses the same shortcuts as the file viewer less, which provides the following actions: Pressing Q will exit the viewer (quit). To scroll up or down half a page press Ctrl+U or Ctrl+D (or Space/Shift+Space for whole page scrolling). /keyword starts a search for "keyword" (only downwards from current position). Pressing N or Shift+N will get you to the next or previous occurence of the keyword. G or Shift+G will goto the beginning or end of the file. To learn more about the manual pager, issue man man.

D.2 Viewing a File

Use 1s to list the files in the current directory (cd .. and cd DIRECTORY to change the directory). For viewing a file there are several possibilities, less FILENAME being one of the more convenient ones (see Section D.1 above for navigation shortcuts). If you want to read a file continuously (because you expect data to be written to it by another program while reading it) you can use tail -f FILENAME, which you have to exit by pressing Ctrl+C.

D.3 Editing a File

There are several file editors you can use on the command line, but if you are not familiar with vim or emacs, the best choice is probably to use nano. You open a file for editing by calling nano FILENAME. Then you can edit the file and move around with arrow keys. All available commands are displayed at the bottom, where ^ denotes the Ctrl key, so you can save the file ("WriteOut") with Ctrl+O and exit nano with Ctrl+X.

D.4 Background Jobs

You can use the & character after any command that may take more time, to start it running in background (so that you can continue to use the terminal while the command is executed).

For example nmap -sP 10.0.0.0/20 > scanresults.txt &

Alternatively, you can press Ctrl+Z (suspend), which stops the current job and sends it to the background. To see all jobs were sent to the background, run jobs which also prints out a number for each job. bg JOBNUMBER or fg JOBNUMBER can be used to either continue a job in the background or foreground. Pressing Ctrl+C while a job is running in foreground will kill (exit) it.

D.5 Clipboard - Copy+Paste

Note that Ctrl+C on the command line is used to kill the current program. Most terminals support Ctrl+Shift+C (after selecting lines with the mouse) and Ctrl+Shift+V as shortcuts for copy and paste to and from the system clipboard. Using the middle mouse button to paste selected text is also supported frequently.